

# Scientific Computing II - Exercise 2

Professor: Antti Kuronen  
Student: Caike Crepaldi

September 2015

This exercise was done using a ssh remote login through my pangolin (pangolin.it.helsinki.fi) university linux account while using my macbook pro with OS X.

To do such a thing, the following command was run in the OS X Yosemite terminal.

Macbook-Pro:~crepaldi \$ ssh username@pangolin.it.helsinki.fi

The editor I used to edit files and write codes was **emacs**.

## 1 First Problem

The Fortran source code of the program created for the first (second and third) problem is shown bellow.

array.f90

```
1 program array
2   implicit none
3   integer :: i, it, a(10)
4
5   write(6,*) "Please give the number iterations :"
6   read(*,*) it
7
8   do i=1,it
9     print *, i, a(i)
10  end do
11
12 end program array
```

The program should ask for the number of desired iterations and print the value of each vector line corresponding to the iteration counter. If the given number of iterations is high enough the program will crash and print the value of each line of the vector until the segmentation fault occurs.

The commands for the compilation and an example of how the program works is shown bellow.

username@tktl-pangolin:~/Tools-for-HPC/exercise-2\$ gfortran -O0 -g -o ex2p1 array.f90

username@tktl-pangolin:~/Tools-for-HPC/exercise-2\$ ./ex2p1

Please give the number iterations:

2000

(The output of the program's execution is really big so I will only show the last 15 lines)

```
941 1966042741
942 796026227
943 926102065
944 3225393
945 791559519
946 1882355813
947 1280245809
948 1146572868
949 1869098813
950 1664050541
951 1885696611
952 6581345
953 2019897134
954 3239986
955 0
956 0
```

Program received signal SIGSEGV: Segmentation fault - invalid memory reference.

Backtrace for this error:

```
#0 0x7FA4B9F9B777
#1 0x7FA4B9F9BD7E
#2 0x7FA4B9BF3D3F
#3 0x7FA4BA064280
```

```
#4 0x7FA4BA065820
#5 0x7FA4BA0694AE
#6 0x400A6E in array at array.f90:9
```

Segmentation fault

We can conclude that the program crashes after trying to access the value of the 957th line of the vector **a**. However, every time the program is run in the linux server, the program crashes while trying to access a different line of the vector **a**.

By running the program 2 more times in a row, we verify that the program crashes first in the 1136th line of the vector and then it crashes in the 861th line of the vector.

## 2 Second Problem

See bellow the commands for such task.

```
username@tktl-pangolin:~/Tools-for-HPC/exercise-2$ gdb ex2p1
(gdb) run
ex2p1
Please give the number interactions:
2000
      1      6311808
      2          0
      3          8
      4          0
      5          0
      6          0
      7          0
      8          0
      9      -5920
     10      32767
     11    4196384
```

(A lot of values will be skiped here)

```
466    791770415
467    1667444579
468    1634755954
469    1278174316
470    2020961897
471    1869567023
472    1714254700
473    1210937967
474    1697596240
475    1668441464
476    761623401
477    2019897138
478    3239986
479          0
480          0
```

```
Program received signal SIGSEGV, Segmentation fault.
0x00007fff7ba2280 in ?? () from /usr/lib/x86_64-linux-gnu/libgfortran.so.3
(gdb) p i
No symbol "i" in current context.
(gdb) p a
$1 = ( (0, 1045149306), 1.2904777690891933e-08 )
(gdb) up
#1 0x00007fff7ba3821 in ?? () from /usr/lib/x86_64-linux-gnu/libgfortran.so.3
(gdb) p i
No symbol "i" in current context.
(gdb) up
#2 0x00007fff7ba74af in ?? () from /usr/lib/x86_64-linux-gnu/libgfortran.so.3
(gdb) p i
No symbol "i" in current context.
(gdb) up
#3 0x000000000400a6f in array () at array.f90:9
9      print *, i, a(i)
(gdb) p i
$2 = 481
(gdb) p a
$3 = (6311808, 0, 8, 0, 0, 0, 0, -5920, 32767)
```

We can see above (in the last four lines) the value of the vector **a** and the last index **i** just before the program crashes.

## 3 Third Problem

In order to create a core dump to be investigated by the **gdb** debugger we must use the following commands.

```
username@tktl-pangolin:~/Tools-for-HPC/exercise-2$ ulimit -c 10000
username@tktl-pangolin:~/Tools-for-HPC/exercise-2$ ./ex2p1
```

```
ex2p1
Please give the number iterations:
2000
```

```
1      6311808
2          0
3          8
4          0
5          0
6          0
7          0
8          0
9      -5920
10         32767
11        4196384
```

(a lot of values will be skiped here)

```
540  1146572868
541  1869098813
542  1664050541
543  1885696611
544  6581345
545  2019897134
546  3239986
547  0
548  0
```

```
Program received signal SIGSEGV: Segmentation fault - invalid memory reference.
```

Now we just need to use the **gdb** debugger to analyse the core dump.

```
username@tktl-pangolin:~/Tools-for-HPC/exercise-2$ ls
array.f90 core ex2p1
username@tktl-pangolin:~/Tools-for-HPC/exercise-2$ gdb ex2p1 -c core

Core was generated by `./ex2p1'.
Program terminated with signal SIGSEGV, Segmentation fault.
#0 0x00007f76ec5bc280 in ?? () from /usr/lib/x86_64-linux-gnu/libgfortran.so.3
(gdb) where
#0 0x00007f76ec5bc280 in ?? () from /usr/lib/x86_64-linux-gnu/libgfortran.so.3
#1 0x00007f76ec5bd821 in ?? () from /usr/lib/x86_64-linux-gnu/libgfortran.so.3
#2 0x00007f76ec5c14af in ?? () from /usr/lib/x86_64-linux-gnu/libgfortran.so.3
#3 0x0000000000400a6f in array () at array.f90:9
#4 0x0000000000400ad9 in main (argc=1, argv=0x7ffcff34fca0) at array.f90:12
#5 0x00007f76ec136ec5 in __libc_start_main (main=0x400aa5 <main>, argc=1,
    argv=0x7ffcff34f8b8, init=<optimised out>, fini=<optimised out>,
    rtld_fini=<optimised out>, stack_end=0x7ffcff34f8a8) at libc-start.c:287
#6 0x0000000000400849 in _start ()
```

## 4 Fourth Problem

The program code can be seen bellow.

simpleprofiling.f90

```
1 program simpleprofiling
2   integer , parameter :: rk=selected_real_kind(5,100)
3   integer , parameter :: NMAX=100000,MMAX=1000
4   real(rk), parameter :: dx=0.00001
5   real(rk) :: sum1,sum2
6   integer :: i,j
7
8   sum1=0.0
9   sum2=0.0
10
11  do j=0,MMAX
12    do i=0,NMAX
13      sum1=sum1+dx*i*j
14    end do
15  end do
16
17  do j=0,MMAX
18    do i=0,NMAX
19      sum2=sum2+cos(exp(sin(dx*i*j)))
20    end do
21  end do
```

```

22
23 print *,sum1,sum2
24
25 end program simpleprofiling

```

According to the assignment, we must look for the time spent in the lines 13 and 19 of the source code.

In order to perform such task we must use the commands bellow.

```

username@tktl-pangolin:~/Tools-for-HPC/exercise-2$ gfortran -O0 -p -g -o simpprof simpleprofiling.f90
username@tktl-pangolin:~/Tools-for-HPC/exercise-2$ ./simpprof
25025249617.809803 24135013.748771366
username@tktl-pangolin:~/Tools-for-HPC/exercise-2$ gprof -l simpprof

```

Flat profile:

```

Each sample counts as 0.01 seconds.
% cumulative self      self      total
time   seconds   seconds   calls  Ts/call  Ts/call  name
44.30    0.92    0.92          1      0.00      0.00  simpleprofiling (simpleprofiling.f90:13 @ 400a0f)
41.87    1.78    0.87          1      0.00      0.00  simpleprofiling (simpleprofiling.f90:19 @ 400a8d)
 6.09    1.91    0.13          1      0.00      0.00  simpleprofiling (simpleprofiling.f90:18 @ 400b00)
 5.36    2.02    0.11          1      0.00      0.00  simpleprofiling (simpleprofiling.f90:12 @ 400a37)
 0.97    2.04    0.02          1      0.00      0.00  simpleprofiling (simpleprofiling.f90:18 @ 400a79)
 0.00    2.04    0.00          1      0.00      0.00  MAIN__ (simpleprofiling.f90:1 @ 4009cd)

%
the percentage of the total running time of the
time   program used by this function.

cumulative a running sum of the number of seconds accounted
seconds   for by this function and those listed above it.

self      the number of seconds accounted for by this
seconds   function alone. This is the major sort for this
listing.

calls     the number of times this function was invoked, if
this function is profiled, else blank.

self      the average number of milliseconds spent in this
ms/call   function per call, if this function is profiled,
else blank.

total     the average number of milliseconds spent in this
ms/call   function and its descendants per call, if this
function is profiled, else blank.

name      the name of the function. This is the minor sort
for this listing. The index shows the location of
the function in the gprof listing. If the index is
in parenthesis it shows where it would appear in
the gprof listing if it were to be printed.

```

Copyright (C) 2012 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification, are permitted in any medium without royalty provided the copyright notice and this notice are preserved.

Call graph (explanation follows)

granularity: each sample hit covers 2 byte(s) for 0.49% of 2.04 seconds

index	% time	self	children	called	name
[6]	0.0	0.00	0.00	1/1	main (simpleprofiling.f90:25 @ 400bb3) [11]
					MAIN__ (simpleprofiling.f90:1 @ 4009cd) [6]

This table describes the call tree of the program, and was sorted by the total amount of time spent in each function and its children.

Each entry in this table consists of several lines. The line with the index number at the left hand margin lists the current function. The lines above it list the functions that called this function, and the lines below it list the functions this one called.

This line lists:

index A unique number given to each element of the table.

Index numbers are sorted numerically.

The index number is printed next to every function name so

it is easier to look up where the function is in the table.

% time This is the percentage of the ‘total’ time that was spent in this function and its children. Note that due to different viewpoints, functions excluded by options, etc, these numbers will NOT add up to 100%.

self This is the total amount of time spent in this function.

children This is the total amount of time propagated into this function by its children.

called This is the number of times the function was called. If the function called itself recursively, the number only includes non-recursive calls, and is followed by a ‘+’ and the number of recursive calls.

name The name of the current function. The index number is printed after it. If the function is a member of a cycle, the cycle number is printed between the function’s name and the index number.

For the function’s parents, the fields have the following meanings:

self This is the amount of time that was propagated directly from the function into this parent.

children This is the amount of time that was propagated from the function’s children into this parent.

called This is the number of times this parent called the function ‘/’ the total number of times the function was called. Recursive calls to the function are not included in the number after the ‘/’.

name This is the name of the parent. The parent’s index number is printed after it. If the parent is a member of a cycle, the cycle number is printed between the name and the index number.

If the parents of the function cannot be determined, the word ‘<spontaneous>’ is printed in the ‘name’ field, and all the other fields are blank.

For the function’s children, the fields have the following meanings:

self This is the amount of time that was propagated directly from the child into the function.

children This is the amount of time that was propagated from the child’s children to the function.

called This is the number of times the function called this child ‘/’ the total number of times the child was called. Recursive calls by the child are not listed in the number after the ‘/’.

name This is the name of the child. The child’s index number is printed after it. If the child is a member of a cycle, the cycle number is printed between the name and the index number.

If there are any cycles (circles) in the call graph, there is an entry for the cycle-as-a-whole. This entry shows who called the cycle (as parents) and the members of the cycle (as children.) The ‘+’ recursive calls entry shows the number of function calls that were internal to the cycle, and the calls entry for each member shows, for that member, how many times it was called from other members of the cycle.

Copyright (C) 2012 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification, are permitted in any medium without royalty provided the copyright notice and this notice are preserved.

Index by function name

```
[6] MAIN__ (simpleprofiling.f90:1 @ 4009cd) [4] simpleprofiling (simpleprofiling.f90:12 @ 400a37) [2]
    simpleprofiling (simpleprofiling.f90:19 @ 400a8d)
[1] simpleprofiling (simpleprofiling.f90:13 @ 400a0f) [5] simpleprofiling (simpleprofiling.f90:18 @ 400
    a79) [3] simpleprofiling (simpleprofiling.f90:18 @ 400b00)
```

We can see in the “*Flat profile*” the time consumption in both lines of code (13 and 19) and it is clear that the time spent in both lines is far greater than the time spent in other lines of the source code.

In the 13th line, the program spent  $\approx 44\%$  of the execution time (0.92 seconds). In the 19th line, the program spent  $\approx 42\%$  of the execution time (0.87 seconds).

If our intention was to optimize the code, we should concentrate in those two lines where most of the execution time is spent. The reason why the time spent in those two lines is much bigger should be the two **do** loops used in each sum calculation.

## Note: About this exercise

All files (and Makefiles, when applicable) created by the student and used in this exercise were packed together with this PDF documentation and shall be available to the assistant for further analysis in its own problem folder.